

DE LA RECHERCHE À L'INDUSTRIE



INTRODUCING KERNEL-LEVEL PAGE REUSE FOR HIGH PERFORMANCE COMPUTING

MSPC 2013

Sébastien Valat¹, Marc Pérache^{1,2}, William Jalby²

¹ CEA,DAM,DIF, F-91297 Arpajon France

² University of Versailles Saint-Quentin en Yveline
45 Avenue des États-Unis, Versailles

CONTEXT : HPC AND MULTI-CORE

- HPC today, **massively parallel**
- Tera 100 : **140K cores** (2010, 1,05 PFlops)
- Nodes are now **multi-core**
- Tera100 / Curie large nodes : **128 cores**
- Nodes have **NUMA** hierarchy
- To exploit such computer we need **MPI + threads**.
- MPC, a unified runtime for **manycore** and **NUMA** architectures.
(MPI 1.3 / OpenMP 2.5 / Pthreads)
- MPC provides a parallel + NUMA memory allocator.



HOW TO MEASURE MALLOC TIME

- How to measure malloc performance :

```
T0 = clock_start();
ptr = malloc(SIZE);
T1 = clock_end();
```

- Ok for **small blocks**, but not for **large** one (> ~128K) :

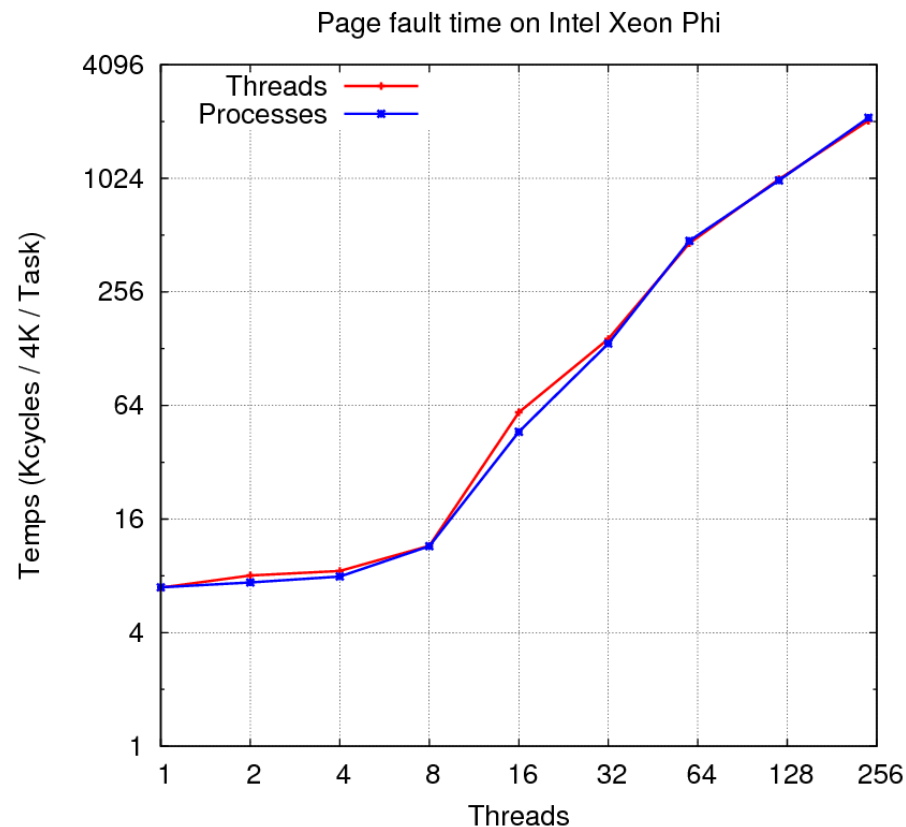

```
T0 = clock_start();
ptr = malloc(SIZE);
for ( i = 0 ; i < SIZE ; i+= PAGE_SIZE)
    ptr[i] = 0;
T1 = clock_end();
```

- Lazy page allocation.

For 4GB	Malloc	First access	Second access
Time (M cycles)	0,008	1 217	5.4

PAGE FAULT SCALABILITY (LINUX 2.6.3X)

- Are page faults scalable over **threads** and **processes** ?
- **Ideally** fault time must be **constant**.
- Measurement on **4*4 Nehalem-EP** (128 cores) and on **Intel Xeon Phi** :



- Page faults are **not scalable** over **threads**

- Some **applications** are **memory intensive** : Hera
 - Large MPI C++ hydrodynamic platform
 - 3D **AMR** meshes
 - Multi-physic / multi-material


- Solutions for applications :
 - Improve applications (not trivial for large one)
 - User-space memory pools (increase memory consumption)
 - **Improve the OS**

COMPARE ALLOCATORS


- **Compare** to production grade **allocators**
- The default one from **glibc**
- Jemalloc (FreeBSD) :
 - Parallel
 - **Lower memory footprint**
 - May generate **too much call to the OS**
- TCMalloc (Google) :
 - Parallel
 - **Keep memory for fast reuse**
 - Get **larger memory consumption**
- MPC :
 - Parallel
 - **Reuse** large memory segments (>1MB)
 - Explicit **NUMA** support
 - Two memory **profiles** (resp. comparable to Jemalloc/TCMalloc)

HERA + MPC ON A BI-WESTMERE (12 : 2 * 6 CORES)

Execution time (s)




Physical memory (GB)




HERA + MPC ON A NEHALEM-EP (128 : 4*4*8 CORES)

Execution time (s)

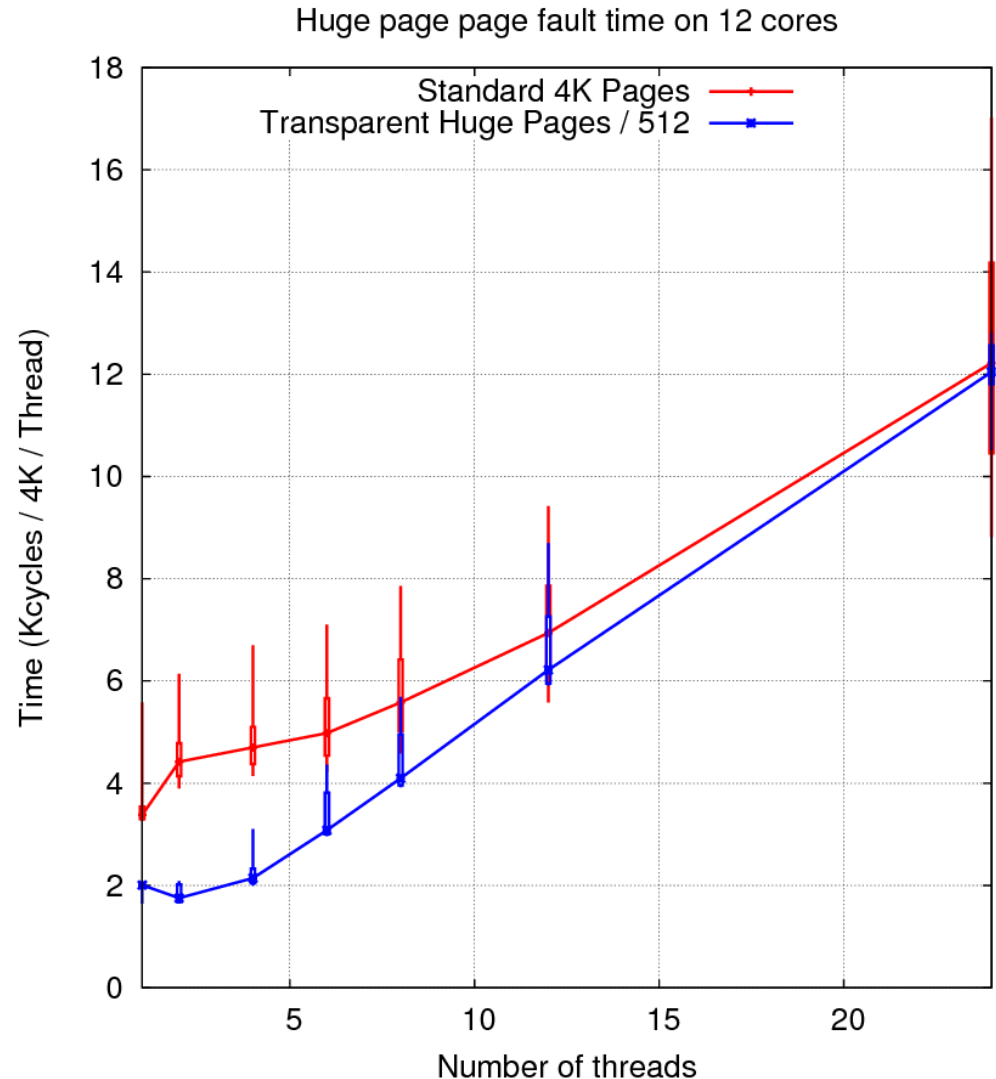


Physical memory (GB)



CAN HUGE PAGES SOLVE THE ISSUE ?

- Standard pages : **4K**
- Huge pages (x86_64) : **2M**
- Divide number of faults by 512
- Can we improve performances ?
 - Sequential : **only 40%**
 - Parallel : **No**
- Why ?



WHAT APPENS ON FIRST TOUCH PAGE FAULT ?

- Hardware generates an interruption
 - Jump to the OS
 - Check reason of the fault
 - Request a free page to NUMA **free lists**
 - **Reset the page content**
 - Map the page, update the **page table**
 - It was done for all 4K pages (262 144 times for 1GB)
- } Possible issue on Xeon Phi
- } ← ~1400/3400 cycles
40%
- } Locks, but hard to fix
some work from
A.T. Clement ASPLOS12

HOW TO AVOID PAGE ZEROING COST

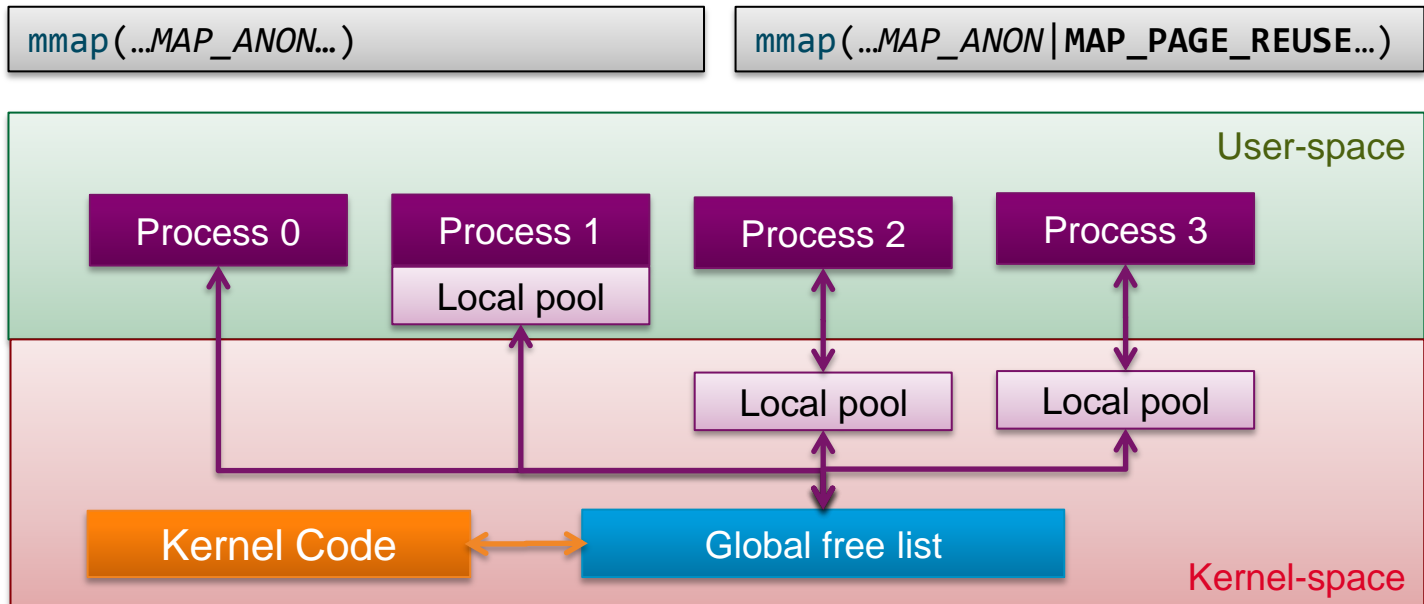
- Windows use a system thread
- So at fault time, pages are already cleaned
- But **zeroing** :
 - Is **unproductive**
 - Consume CPU **cycles** so **energy**
 - Consume **memory bandwidth**
- Why not to **avoid them** ?
- Most allocations pattern follow :

```
double * ptr = malloc(SIZE * sizeof(double));  
for ( i = 0 ; i < SIZE ; i++)  
    ptr[i] = default_value(i);
```

- Why not **inform the kernel** that we do not need zeros ?

REUSING LOCAL PAGES TO AVOID ZEROING

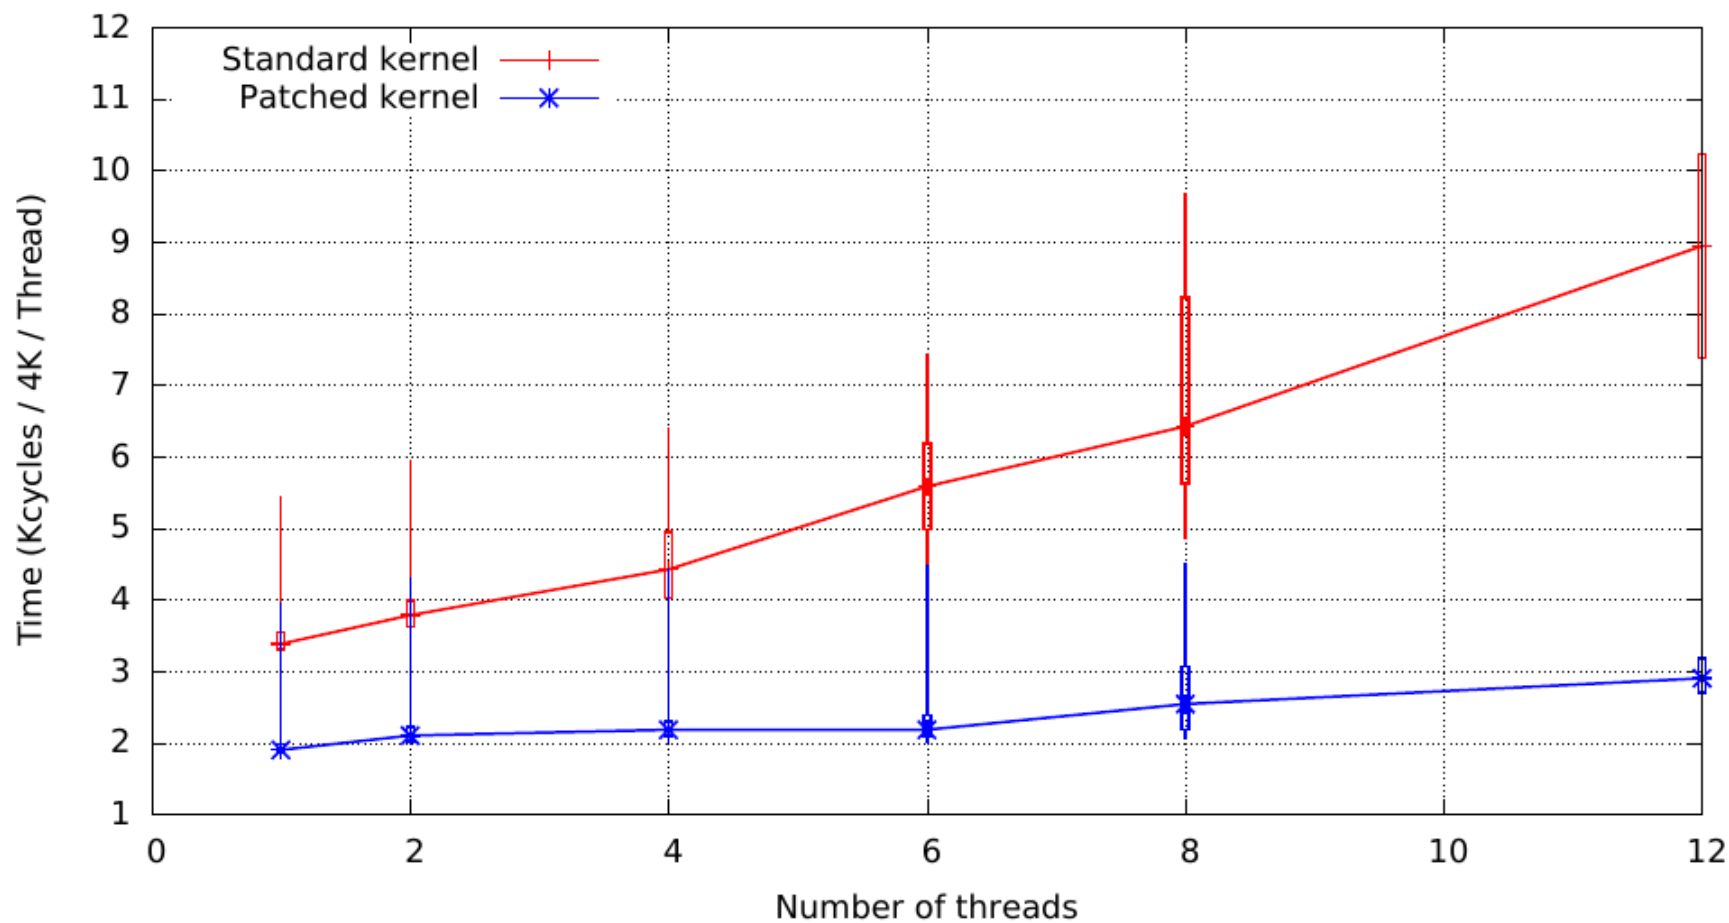
- We can **extend** the **mmap** semantic :
- Page zeroing is **required** for **security reason**
- It prevent information **leaks** from **another processes** or from the **kernel**.
- **But we can reuse pages locally !**



IMPROVEMENT OF FAULTS ON 6 CORE WESTMERE

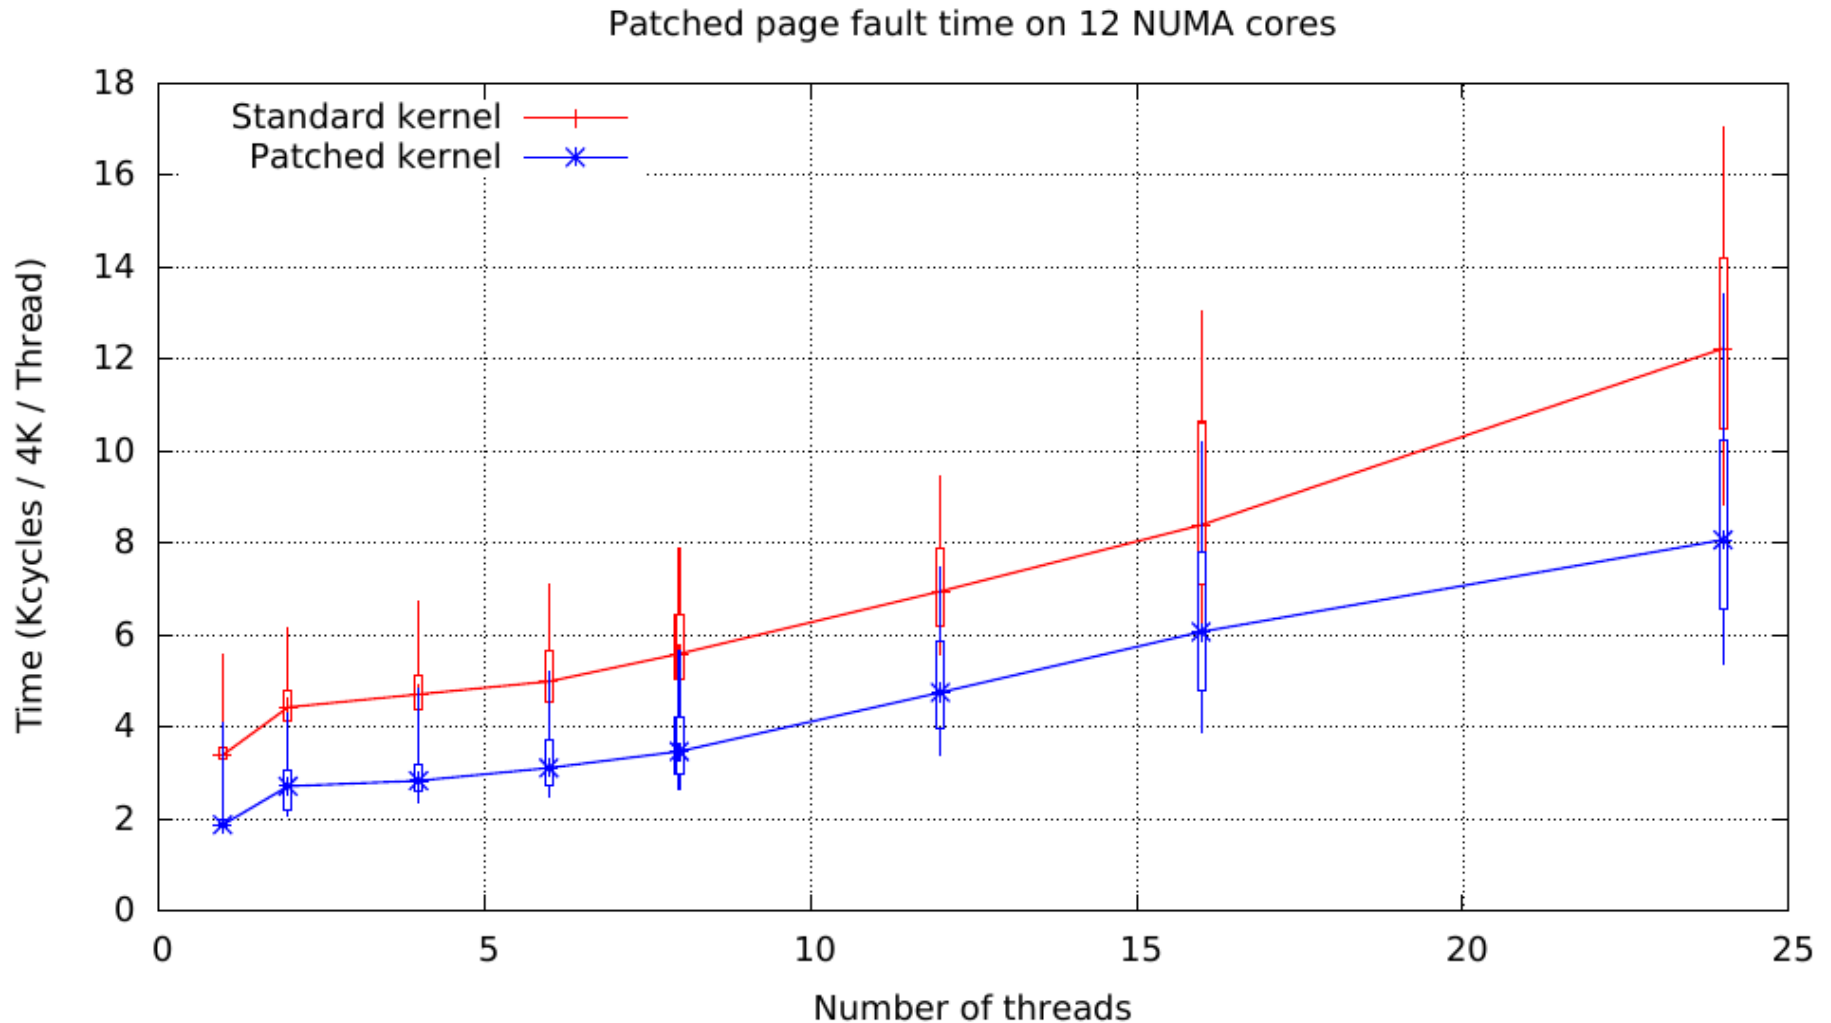
- Without NUMA, get good results.

Patched page fault time on 1 socket of 6 cores



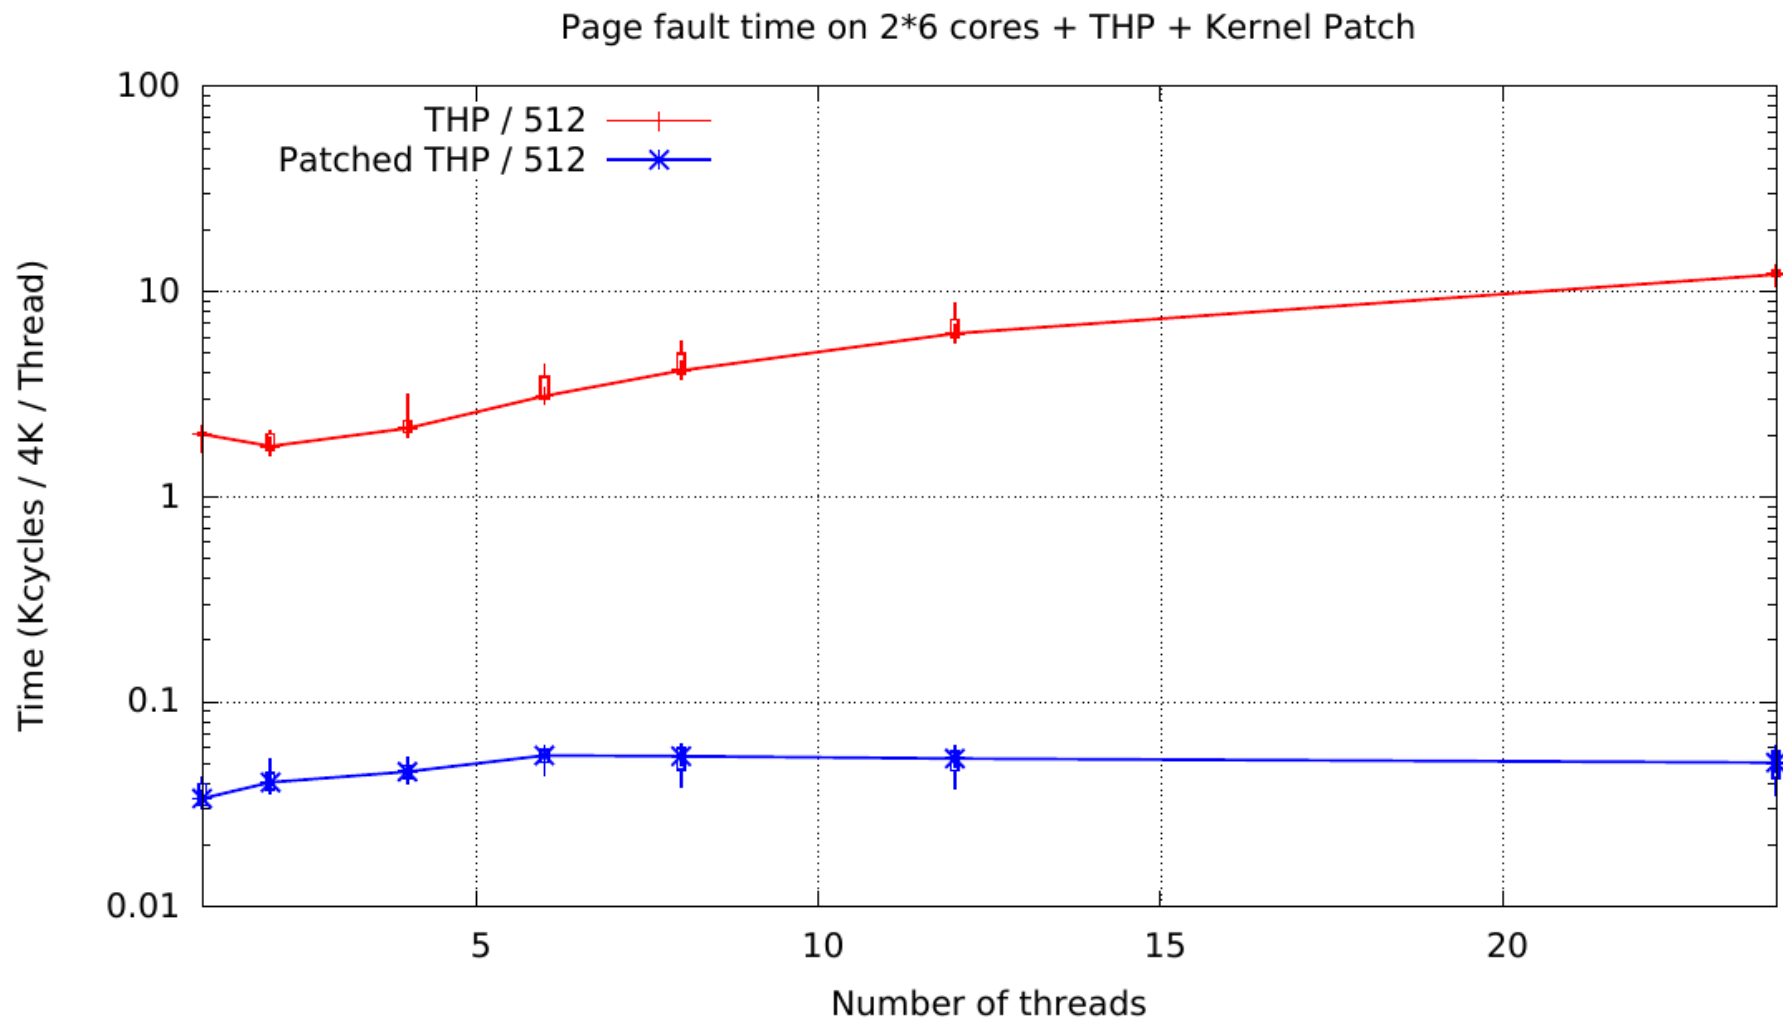
USING TWO SOCKETS (NUMA)

- Our patch improve performance, but **NUMA effects** due to locks became dominant.



EFFECT ON HUGE PAGES

- Get huge improvement (x60), **new interest for huge pages.**



RESULTATS HERA SUR BI-WESTMERE (2*6 COEURS)

■ Standad pages (4K) :

Allocator	Kernel	Total (s)	Sys. (s)	Mem. (GB)
Glibc	Std.	144	9	3,3
MPC-NUMA	Std.	135	2	4,3
MPC-Lowmem	Std.	162	16	2,0
MPC-Lowmem	Patched	157	11	2,0
Jemalloc	Std.	143	15	1,9
Jemalloc	Patched	140	9	3,2

■ Transparent Huge Pages (2M) :

Allocator	Kernel	Total (s)	Sys. (s)	Mem. (GB)
Glibc	Std.	150	13	4,5
MPC-NUMA	Std.	138	2	6,2
MPC-Lowmem	Std.	196	28	3,9
MPC-Lowmem	Patched	138	3	3,8
Jemalloc	Std.	145	15	2,5
Jemalloc	Patched	138	6	3,2

WHERE TO PLACE MEMORY POOLS ?

	User-space	Kernel-space	
Sizes	*	4KB	2MB
Controlling memory	Virtual	Physical	
Limit mono process consumption	~	+	
Limit multi-processes consumption	-	~ / + ?	
Adaptation to real access pattern	-	+	
Ease of implementation	+	-	
Support of NUMA	~	+	
Performance gain	+	~	+

Conclusion

- **Page zeroing** account for **~40%** in sequential !
- **Extend** mmap/madvise **semantic** to remove need of page zeroing.
- Get the expected 40% sequential improvement with 4K pages.
- **New interest for huge pages** (reduction of x60).

Future work and open question

- Integration in page reclaim algorithm.
- Still limited by **lock scalability** on **NUMA** !
- What is a **good huge page size** ? 2M too large ?

QUESTIONS ?

BACKUPS

- Need to find **balance** between **consumption** vs. **performances**.
- On **128 cores**, improvement of **20%** for **2GB**.
- With **NUMA** support, improvement of **48%** compared to the best one.
- Can we **improve** the **consumption / performance ratio** ?

INTEGRATION WITH PAGE RECLAIM

- **Consumption** currently limited to the **maximum working set** of the application.
- Need more work to support “**page reclaim**” in case of memory famine.
- Page reclaim functions need to loops overs local pools before swaping.
- In case of repetitive reclaims, disable usage of local pools until lower memory pressure.
- General aspects of swap integration was looked, but not implemented.

- Codes can rely on lazy page zeroing!
- Cannot enable it by default.
- Need explicit demand with mmap / madvise flags :

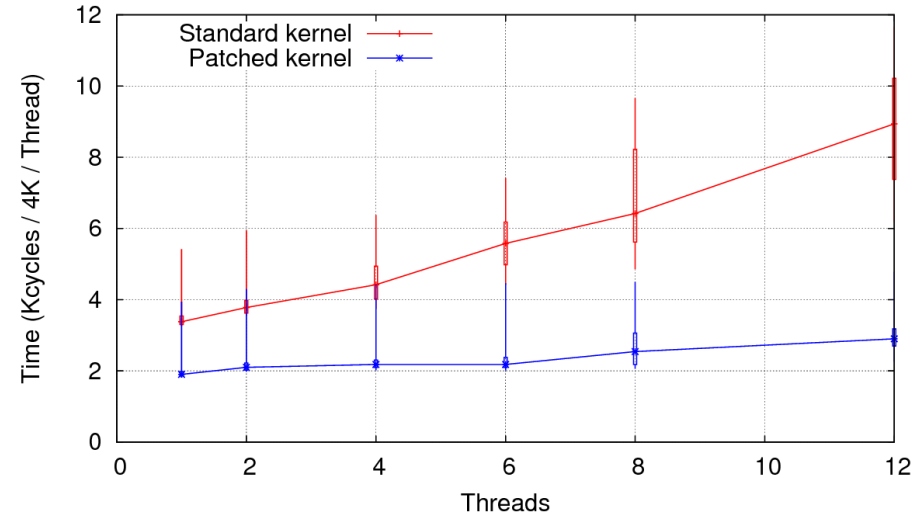
```
void * ptr = mmap(... MAP_ANON | MAP_PAGE_REUSE ...);  
munmap(ptr);  
  
ptr = mmap(... MAP_ANON ...);  
madvise(...MADV_PAGE_REUSE);  
munmap(ptr);
```

- Need to patch malloc/realloc, tested support in
 - MPC_Allocator
 - Jemalloc

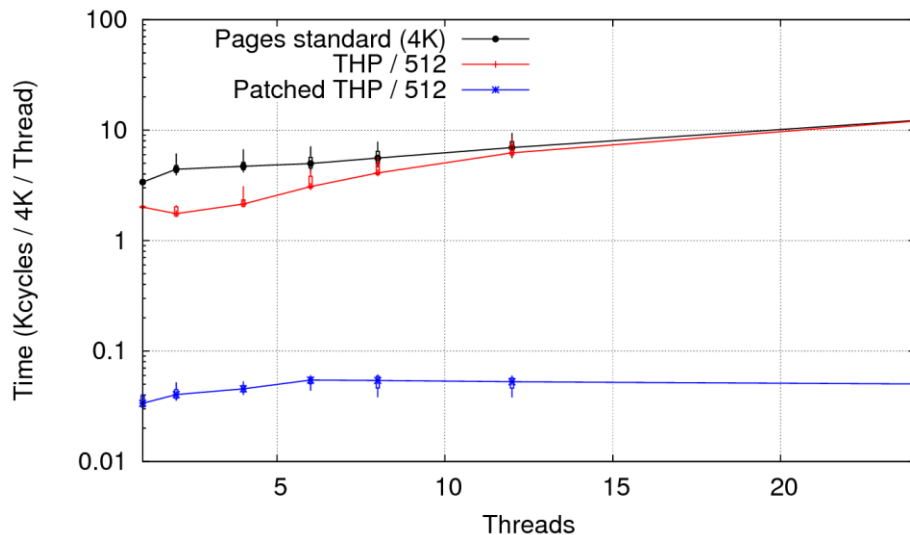
IMPROVEMENT OF FAULTS ON BI-WESTMERE

- Improvement for 4K pages
- But dominated by **NUMA effects** on kernel locks
- Large impact on **huge pages** (2M)
- Observe **new interest** for huge pages


Patched page fault time on 1 socket (6 cores)



Patched page fault time on 2*6 cœurs + THP



Patched page fault time on 2 socket (12 cores)



INTEGRATION WITH PAGE RECLAIM

- Current implementation implicitly limits the memory consumption to the maximum working set of the application.
- Need more work to support “page reclaim” in case of memory famine.
- Page reclaim functions need to loops overs those local caches before swaping.
- In case of repetitive reclaims, disable usage of local pools until lower memory pressure.
- General aspects of swap integration was looked, but not implemented.

HYDRO RESULTS ON BI-WESTMERE (2*6 CORES)

■ Kernel patch and standard 4K pages

App.	Allocator	Kernel	Total (s)	Sys. (s)	MFlops
Std.	Glibc	Std.	1:29	30,7	1770
Std.	MPC	Std.	1:28	31,5	1775
Std.	MPC	Patched	1:19	19,7	2004
Std.	MPC-KeepMem	Std.	0:59	0,5	2649
Patch.	Glibc	Std.	0:43	0,4	3606

■ Kernel patch and Transparent Huge Pages

App.	Allocator	Kernel	Total (s)	Sys. (s)	MFlops
Std.	Glibc	Std.	1:13	18,8	2140
Std.	MPC	Std.	1:18	17,8	2007
Std.	MPC	Patched	1:11	7,0	2224
Std.	MPC-KeepMem	Std.	1:05	1,0	2412
Patch.	Glibc	Std.	0:50	0,4	3554

- Some applications are **intensive** in **memory allocations**

- Application **Hera**:
 - Large MPI C++ hydrodynamic platform
 - 3D **AMR** meshes
 - Multi-physic / multi-material
 - We used it with MPC thread-based MPI.

- Application HydroBench:
 - A smaller hydrodynamic MPI / OpenMP benchmark
 - An older version generate large number of memory allocations.


- Memory management can become a **bottleneck**

- OS (Linux) memory management scalability ?


- Need to find **balance** between **consumption** vs. **performances**.
- On **128 cores**, improvement of **20%** for **2GB**.
- With **NUMA** support, improvement of **48%** compared to the best one.
- Can we **improve** the **consumption / performance ratio** ?

IMPROVEMENT OF FAULTS ON 6 CORE WESTMERE

1 thread



12 threads (hyper-threading)

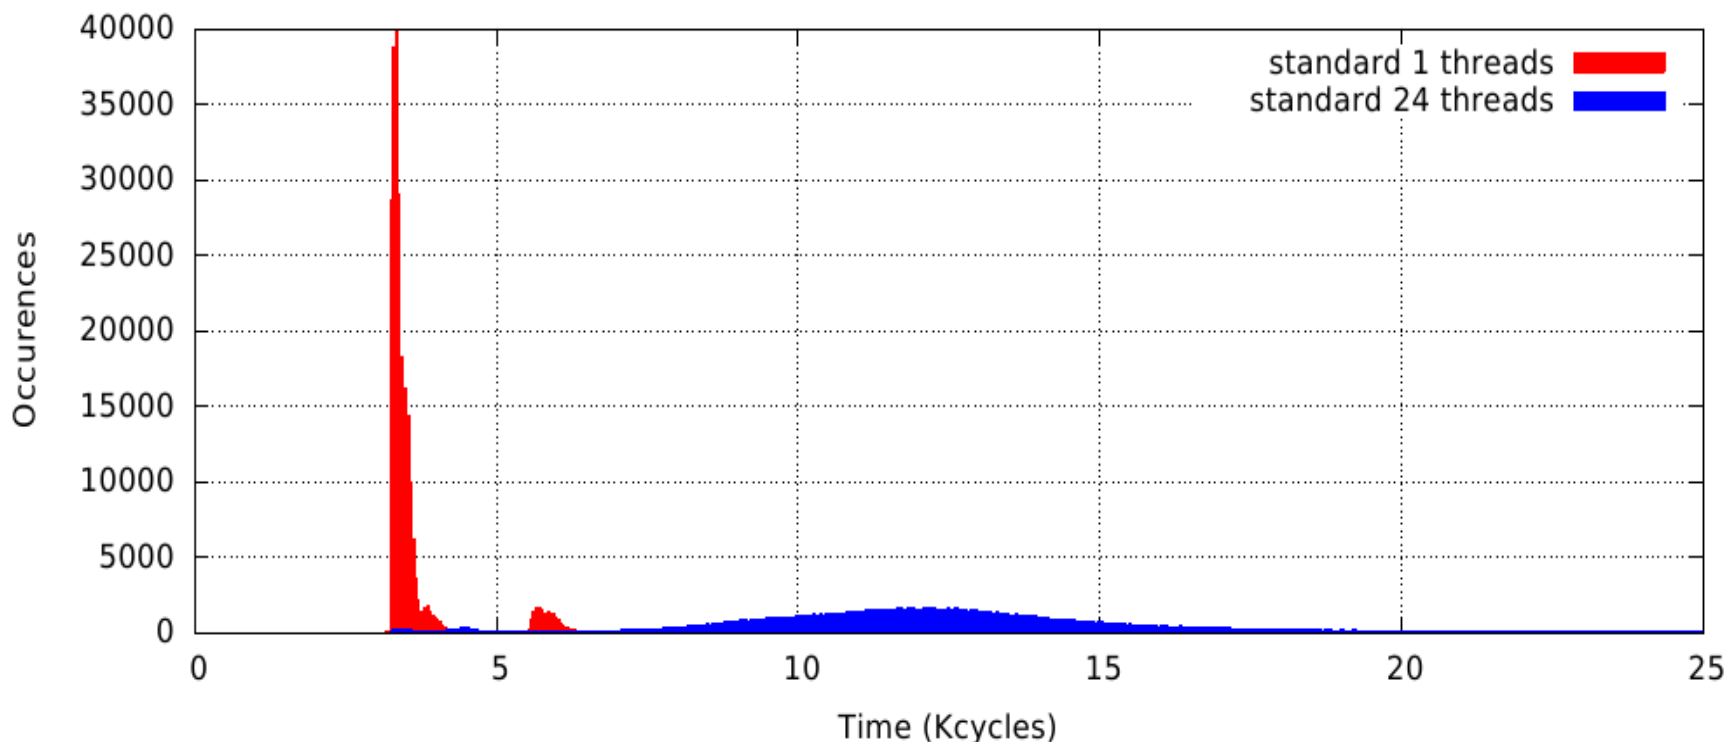


- A **unified runtime** to support **MPI + X** applications.
- Implement **standards** :
 - **MPI 1.3**
 - **OpenMP 2.5**
 - **Pthread**
- Optimized for **manycore** and **NUMA** architectures.
- Provide a **thread-base MPI** mode (tasks are threads, not processes).
- **Be interested in thread performances.**
- **Need parallel and NUMA aware memory allocator.**

MEASURE PAGE FAULT TIME DISTRIBUTION


- Measure each fault with RDTSC (first access to fresh memory segments).
- Obtain time distribution by repeating many times to observe variability.
- In sequential or parallel.

Page fault time distribution (per 4K pages per thread)




HERA + MPC ON A NEHALEM-EP (32 : 4 * 8 CORES)

Execution time (s)



Physical memory (GB)



SOME DETAILS

Reuse policy :

- Search the best fitting segment.
- Rely on mremap to reuse segments which do not fit with the request.

Limit the consumption :


- Keep all segments smaller than configured size (~20 MB).
- Limit the total amount of unused memory (~4GB per NUMA node).

Limitations

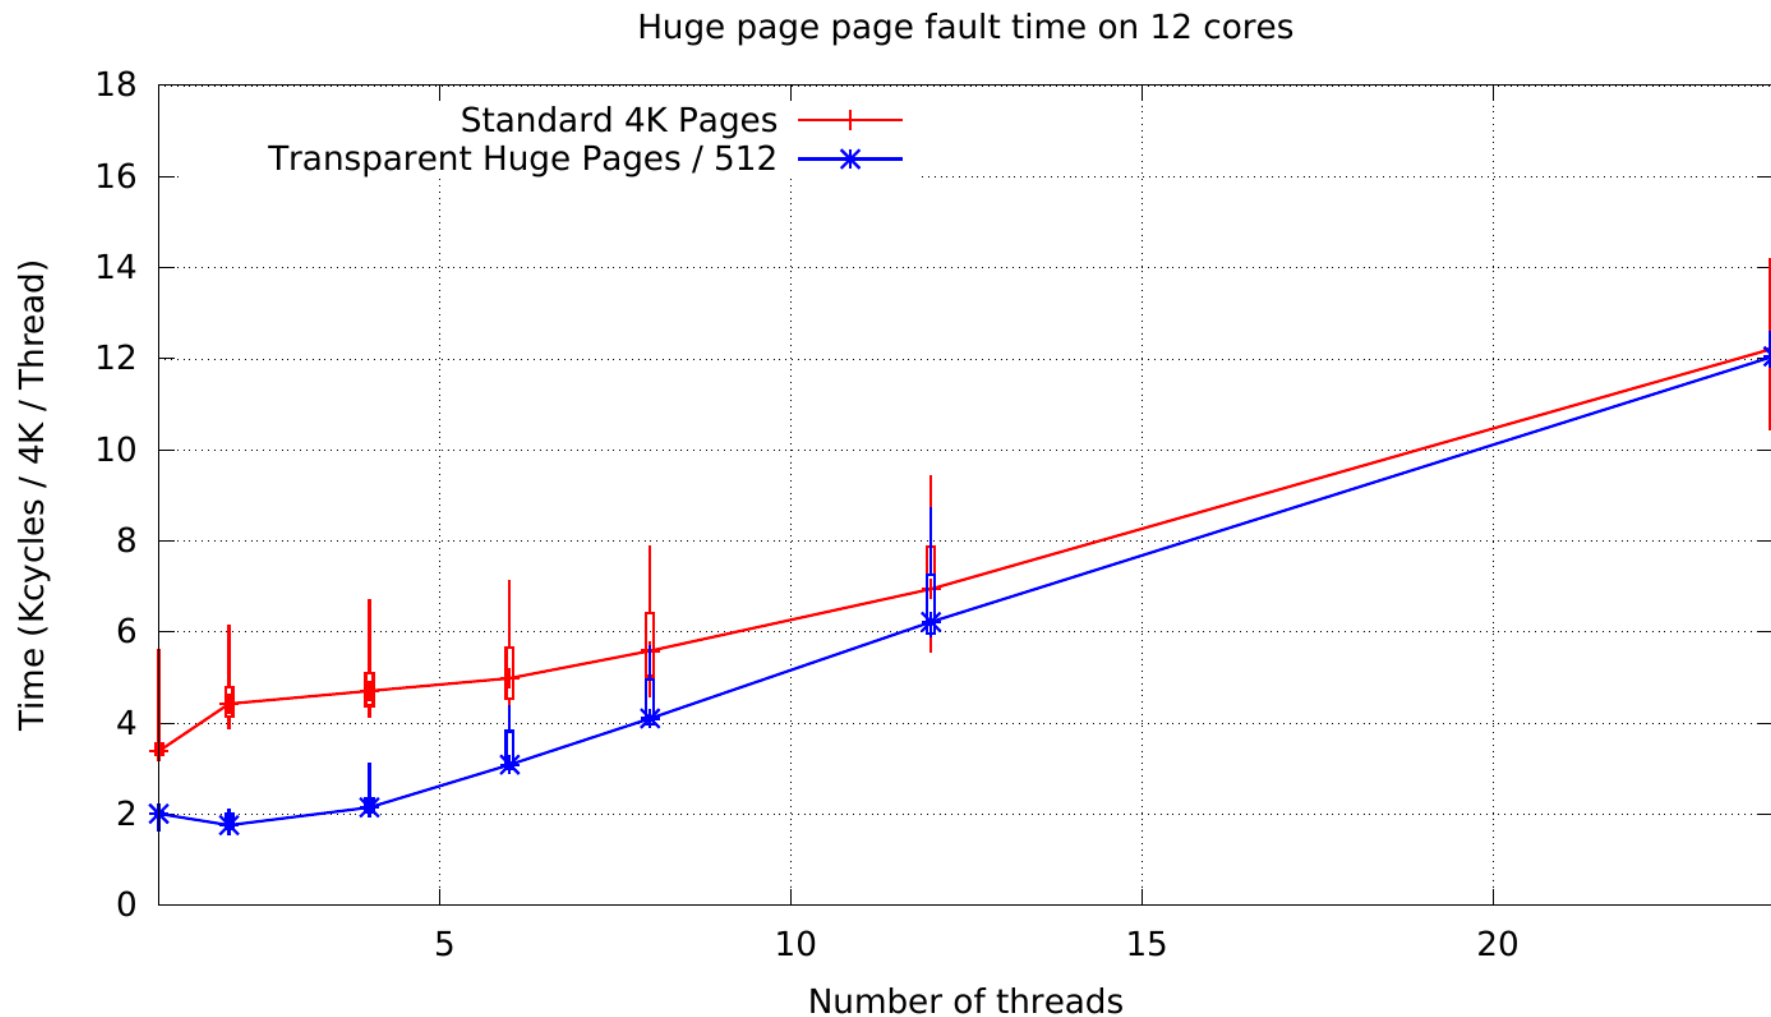
- Larger memory consumption.
- Side effects with application which allocate more memory than required.
- Application dependent parameters, need to automate.

ALLOCATION AND PAGE FAULT SEMANTIC

- mmap reserves memory regions (in Linux kernel, VMA : Virtual Memory Area).
- Regions are initially not provisioned in pages.
- Pages are really mapped on first access (first touch).




HUGE PAGES SCALABILITY




IMPROVEMENT OF FAULTS ON 6 CORE WESTMERE

1 thread




12 threads (hyper-threading)

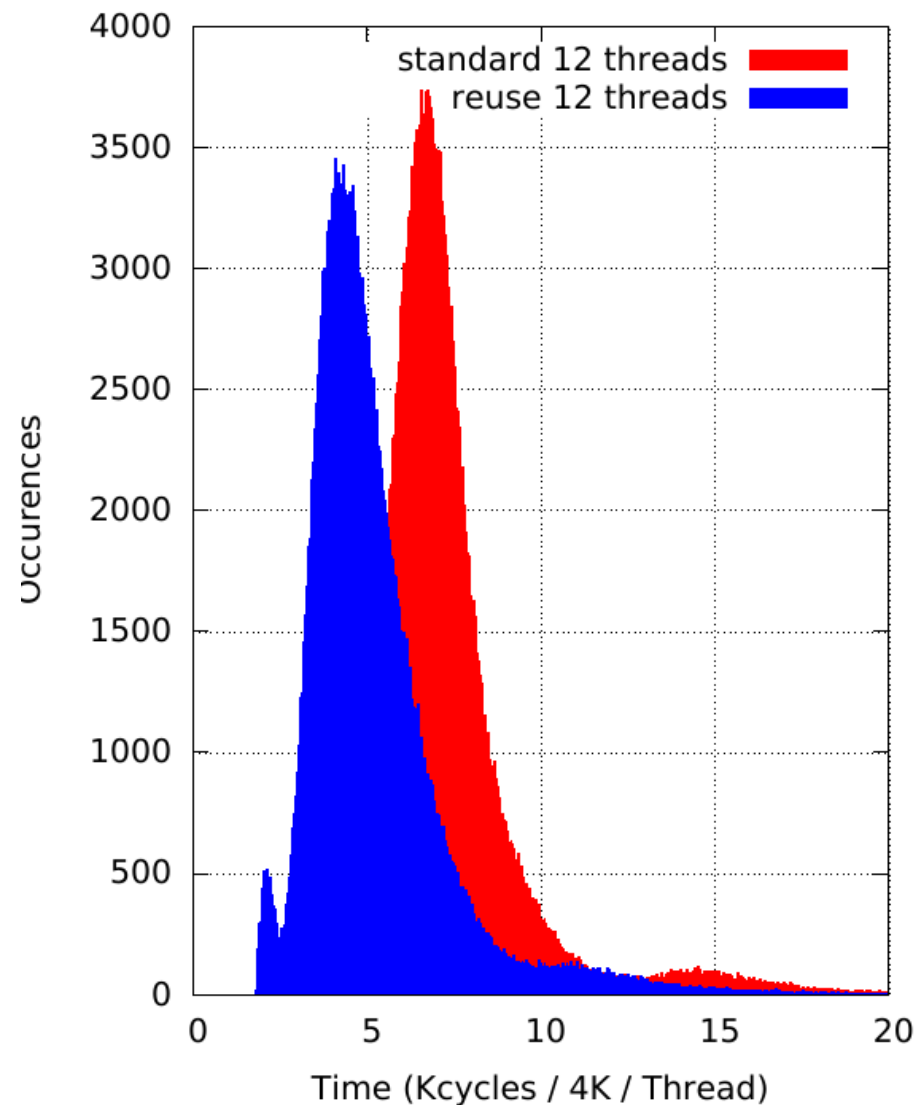


USING TWO SOCKETS (NUMA)

One socket (UMA)




Two sockets (NUMA)



REDUCE MEMORY EXCHANGES WITH THE OS

- Introduce a **pool between application and OS**
- Reuse large segments in memory allocator.
- Require **explicit NUMA** support in whole allocation chain to be efficient.
- **Current allocators** do reuse for small segments, but **not for large** (> 1MB).



- Page faults are **not scalable** over **threads**
- Improve applications ? (not trivial for large one)
- User-space memory pools ? (increase memory consumption)
- Improve the OS ?

CLEAR PAGE COSTS

- Page fault cost in mean : ~**3400** cycles
- Clear page cost in mean : ~**1400** cycles
- On page fault, **40%** of the time is due to **zero filling** !
- Clear page function is called **between** two read **locks**
- It prevent parallel usage of mmap/munmap/brk.
- On huge pages it has to clear **2MB** instead of **4KB** per page fault.